

Master Thesis: *Natural gas price forecasting using recurrent neural networks*

Christian Koopmann

School of Business and Economics

Chair of Information Systems

Humboldt-Universität zu Berlin

<http://https://www.wiwi.hu-berlin.de/de/professuren/quantitativ/wi>



Outline

Motivation

Natural Gas Market

Methodology

Input Data

Training / Tuning Approach

Price Level Prediction

Binary Prediction

Conclusions / Outlook / Comments

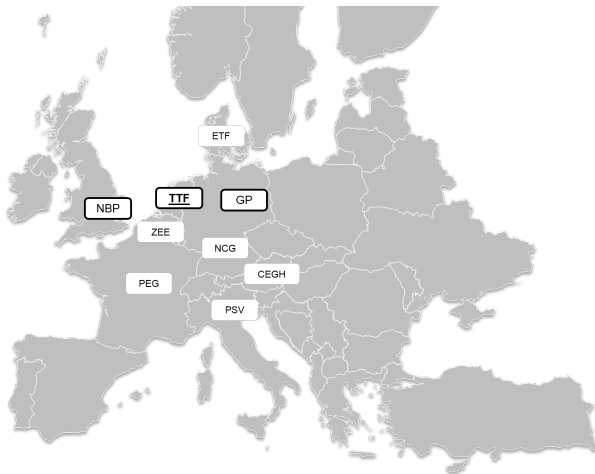
Motivation

- *Methodology:* Evaluate the value of Long Short Term Memory Recurrent Neural Networks for time series prediction
 - ▶ Compare performance to simple RNN, FFNN, and linear reference models
- *Application:* Support natural gas traders in choosing the optimal trading strategy
 - ▶ Trading as part of industrial procurement to meet physical demand
 - ▶ No speculative trading

Description of the target variable / gas prices

- ▣ Different types of natural gas futures differ in various ways
- ▣ The **Virtual Trading Point** describes in which part of the European natural gas transport network the gas is delivered
- ▣ The **Delivery Period** describes during which time frame the gas is delivered at a constant rate
- ▣ The target variable is the future price of gas traded at the **TTF VTP** for delivery in the **next calendar month**.

Virtual Trading Points




Price History TTF Front Month



Why RNN ?

- ▣ FFNN can only learn static input-output mappings
- ▣ For machine learning problems based on sequential data the input-output mapping should be dynamic
- ▣ Examples of sequential data: Text, Speech, Videos, Financial Time Series
- ▣ RNN's are able to learn dependencies of arbitrary length, which does not need to be specified.
- ▣ Main idea: use hidden layer output at one time point as input to the hidden layer at the next input


General Layout



Simple RNN

- Different types of RNN differ in the way they connect the hidden layers between time steps
- The simplest variant treats the previous output h_t in the same way as the other inputs x_t
- $h_t = f(Wh_{t-1} + Ux_t)$

Hidden Layer of Simple RNN



Vanishing Gradient Problem

- Recursive definition of the hidden layer can be expanded:

- ▶ $h_t = f(W h_{t-1} + U x_t)$

- ▶ $h_t = f(W f(W h_{t-2} + U x_{t-1}) + U x_t)$

- Repeated application of chain rule:

- ▶ $\frac{h_t}{dW_{ij}} = \frac{df(z)}{dz} \left(\frac{dW}{dW_{ij}} h_{t-1} + W \frac{dh_{t-1}}{dW_{ij}} \right)$

- ▶ $\frac{h_t}{dW_{ij}} = \frac{df(z)}{dz} \left(\frac{dW}{dW_{ij}} h_{t-1} + W \frac{df(z)}{dz} \left(\frac{dW}{dW_{ij}} h_{t-2} + W \frac{dh_{t-2}}{dW_{ij}} \right) \right)$


Vanishing Gradient Problem

- Contribution of $\frac{dh_{t-k}}{dW_{ij}}$ is multiplied by $(\frac{df(z)}{dz})^k W^{k-1}$
 - ▶ Exponential behaviour leads to either *exploding* or *vanishing gradient problem*
 - ▶ *Exploding* case can be controlled relatively easily by clipping the gradient
 - ▶ No solution of *vanishing gradient problem* in this model set-up
 - ▶ RNNs unable to learn long term dependencies

LSTM

- Long Short Term Memory Networks try to overcome this problem by introducing the cell state c_t
 - ▶ c_t is manipulated in different *gates*
 - ▶ In each gate c_t is multiplied by or added to the output of a layer of neurons with trained weights
- Formal definition:
 - ▶ $h_t = \sigma_o(W_4 h_{t-1} + U_4 x_t) * f_o(c_t)$
 - ▶ $c_t = \sigma_f(W_1 h_{t-1} + U_1 x_t) * c_{t-1} + \sigma_i(W_2 h_{t-1} + U_2 x_t) * f_i(W_3 h_{t-1} + U_3 x_t)$
- LSTM does not suffer from *vanishing gradient problem* but has four times as many parameters to train with same input data.

Hidden Layer of LSTM



Data Overview

- All Data was downloaded from the *Thomson Reuters Eikon* data base.
- **Energy Commodity Prices:** Target variable, other gas futures, Brent Oil Futures, Coal Futures, Electricity Base / Peak futures
- **Exchange Rates:** EUR/GBP, EUR/USD
- **Gas Market Fundamentals:** Storage Levels, Pipeline Flows, National Consumption / Production Data
- All data as daily values (Closing prices) starting between 2010 - 2014.

Training / Tuning Approach

1. Parameter Tuning of univariate models
 - ▶ Single Train / Test split
 - ▶ Training Data 2010 - 2015 / Test Data: 2016
 - ▶ Tuning of: Network Architecture, Dropout, Learning Rate
2. Variable Selection of multivariate models
 - ▶ Use tuned parameters from respective univariate model
 - ▶ Forward variable selection based on MSE
 - ▶ Same Train / Test split as above
3. Parameter Tuning of multivariate models
 - ▶ Use previously selected input variables
 - ▶ Same parameters / data as in univariate case
4. Model evaluation
 - ▶ Month wise cross validation with testing months selected from 01 - 08/2017

Predicting Price Levels

- Predict tomorrows closing price of the TTF Front Month future based on all data available up to the current day
- The MSE loss function is minimized using stochastic gradient descent
- Linear Reference Model: AR(1)
 - ▶ Parameter estimate very close to 1, predictions almost identical to current price value

Price Level Prediction Results

	Model	Variables	MSE	MSEReference
1	lstm		0.143	0.151
2	rnn		0.150	0.151
3	ffnn_long		0.151	0.151
4	ffnn_long	EURGBPFX_EURUSDFX_TradeUK_TradeRUNWE	0.152	0.151
5	rnn	TTFDA	0.154	0.151
6	ffnn_short		0.162	0.151
7	ffnn_short	EURGBPFX_TTFDA_ProdNL_TradeBBL_NBPfM	0.162	0.151
8	lstm	TradeUK	0.181	0.151

Table 1: Test Results using monthly cross validation of tuned models for data 01 - 07/2017

Binary Prediction



□ Problem:

- ▶ Decision Problem: Buy today or at any other day before the end of the month to cover physical demand
- ▶ Once bought, futures can't be sold again (Company Policy, Regulation).
- ▶ Therefore optimal trading strategy can't be derived directly from tomorrows price level

□ Solution:

- ▶ New binary target variable: *Is today's closing price minimal among all closing prices for the rest of the month ?*
- ▶ Yes - 1, No - 0
- ▶ Loss Function: Binary Cross Entropy
- ▶ Naive reference model: $\tilde{y}_i = \frac{1}{N_i}$ with N_i the remaining trading days for this month

Binary Prediction



Binary Prediction Results

	Model	Variables	CrossEntropy	CrossEntropyReference
1	mlp_long	EURUSDFX	0.447	0.497
2	rnn	EURGBPFX	0.454	0.497
3	ffnn_short	ElectricityBaseFM_EURUSDFX	0.454	0.497
4	lstm		0.468	0.497
5	lstm	ProdUKCS	0.469	0.497
6	rnn		0.486	0.497
7	ffnn_short		0.487	0.497
8	ffnn_long		0.517	0.497

Table 2: Test Results using monthly cross validation of tuned models for data 01 - 07/2017

Conclusions

- ▣ Among univariate models LSTM outperforms alternative models in both prediction problems
- ▣ Among multivariate problems opposite seems to be the case
- ▣ Univariate LSTM shows best relative performance in Price Level Prediction where it is the only model to significantly outperform the Linear Reference
- ▣ Univariate models seem to be better in Price Level Prediction where the opposite is true for the binary case.

Conclusions

There are several ways for possible extension / improvement

- ▣ Extend parameter tuning to choice of optimizer, activation, length, batch size etc.
- ▣ Extend parameter tuning to include multi-level architectures
- ▣ Use a finer tuning grid for the current parameters
- ▣ Use Cross Validation during parameter tuning

Comments

- Model Training / Tuning: Python (Keras, Tensorflow), Result Analysis: R
- Running models on Amazon Web Services, is easier / cheaper than expected and frees valuable resources on local machine
- Valuable Resources on LSTMs and implementation in Python:
 - ▶ BlogPost *Understanding LSTM Networks*:
<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>
 - ▶ ML Blog *Machinelearningmastery*:
<https://machinelearningmastery.com/>

Binary Prediction Univariate parameter tuning

	Model	Architecture	Dropout	LearningRate	binary_crossentropy
1	lstm	16	0.250	0.001	0.450
2	rnn	8	0.250	0.001	0.449
3	mlp_long	32	0.000	0.001	0.468
4	mlp_short	8	0.000	0.010	0.473

Table 3: Best parameter combinations for each binary model

Binary Prediction Variable Selection

	Model	Variables	binary_crossentropy
1	lstm	ProdUKCS	0.45
2	rnn	EURGBPFX	0.44
3	mlp_long	EURUSDFX	0.47
4	mlp_short	ElectricityBaseFM_EURUSDFX	0.47

Table 4: Best variable combinations for each binary model

Binary Prediction Multivariate parameter tuning

	Model	Variables	Architecture	Dropout	LR	CE
1	lstm	ProdUKCS	32	0.500	0.001	0.452
2	rnn	EURGBPFX	8	0.250	0.001	0.444
3	mlp_long	EURUSDFX	32	0.000	0.001	0.478
4	mlp_short	ElectricityBaseFM_EURUSDFX	16	0.000	0.010	0.461

Table 5: Best parameter combinations for each binary model

Price Level Prediction Univariate parameter tuning

	Model	Architecture	Dropout	LearningRate	mse
1	lstm	8	0.000	0.001	0.117
2	rnn	8	0.000	0.001	0.122
3	mlp_long	8	0.000	0.100	0.736
4	mlp_short	16	0.000	0.100	0.459

Table 6: Best parameter combinations for each binary model

Price Level Prediction Variable Selection

	Model	Variables	mse
1	lstm	EURUSDFX	0.12
2	rnn	ElectricityBaseFM	0.11
3	mlp_long	TradeRUNWE_ConLDZNL	0.16
4	mlp_short	TTFDA_StorageEU_LNGStockEU_ProdNL	0.22

Table 7: Best variable combinations for each level prediction model

Price Level Prediction Multivariate parameter tuning

	Model	Variables	Architecture	Dropout	LR	ms
1	lstm	EURUSDFX	32	0.250	0.001	0.11
2	rnn	ElectricityBaseFM	8	0.000	0.001	0.14
3	mlp_long	TradeRUNWE_ConLDZNL	16	0.000	0.100	0.18
4	mlp_short	TTFDA_StorageEU_LNGStockEU_ProdNL	8	0.000	0.100	0.17

Table 8: Best parameter combinations for each binary model